

# SFWRENG 4NL3 Assignment 4: Pretrained Transformers

Sumanya Gulati

1 April 2025

Contents

1 Dataset 2

1.1 Data Collection . . . . . 2

1.2 Dataset Structure . . . . . 2

1.3 Evaluation Metrics . . . . . 3

1.4 Data Splits . . . . . 3

1.5 Additional Features . . . . . 4

2 Fine-tuned Models 4

2.1 Model 1: RoBERTa-base . . . . . 5

2.1.1 Pretraining Dataset . . . . . 5

2.1.2 Compute Requirements for Pretraining . . . . . 5

2.2 Model 2: DistilBERT-base-uncased . . . . . 5

2.2.1 Pretraining Dataset . . . . . 5

2.2.2 Compute Requirements for Pretraining . . . . . 5

2.3 Fine-Tuning Steps . . . . . 6

List of Tables

1 Label Mapping . . . . . 3

2 Data Split Overview . . . . . 4

List of Figures

1 Class Distribution in Training and Test Splits . . . . . 4

# 1 Dataset

For this assignment, I am using the GoEmotions dataset extracted from the HuggingFace can be accessed [here](#). The task at hand involves emotion classification, aiming to identify and categorize the emotions expressed in textual data. This is essential for applications such as sentiment analysis, mental health assessment, enhancing human-computer interactions and more.

The GoEmotions dataset comprises of about 58,000 English Reddit comments, each annotated for 27 distinct emotion categories or marked as neutral. The simplified version of the dataset with predefined train, val and test splits has been used for this assignment.

## 1.1 Data Collection

The dataset has been constructed by selecting English-language comments from Reddit by researchers at Amazon Alexa, Google Research and Stanford Linguistics. A complete list of authors can be found [here](#). The comments were extracted from Reddit using a variety of automated methods along with data curation techniques such as reducing profanity, length filtering, sentiment and emotion balancing, masking and more. Further information about the data collection process can be found in section 3.1 of this paper.

## 1.2 Dataset Structure

Each instance of the dataset corresponds to a reddit comment with an ID and one or more emotion annotations including neutral. The simplified configuration of the dataset which has been used for this assignment, includes:

- **text**: the Reddit comment
- **labels**: the emotional annotations
- **comment\_id**: a unique identifier for the comment

The input for the task is a Reddit comment in English and the corresponding output is a set of one or more labels corresponding to the 27 emotion categories or neutral, reflecting the emotional content of the comment.

The labels are stored as a list of integers ranging from 0 to 27 where each integer represents an emotion category or neutral. The label mapping is as follows:

Label Number	Label Category
0	admiration
1	amusement
2	anger
3	annoyance
4	approval
5	caring
6	confusion
7	curiosity
8	desire
9	disappointment
10	disapproval
11	disgust
12	embarrassment
13	excitement
14	fear
15	gratitude
16	grief
17	joy
18	love
19	nervousness
20	optimism
21	pride
22	realization
23	relief
24	remorse
25	sadness
26	surprise
27	neutral

Table 1: Label Mapping

### 1.3 Evaluation Metrics

The following evaluation metrics have been used to assess the performance of the BERT-based model:

- Model Performance:
  - Emotion-level Precision, Recall, F1: Measured per each emotion in the GoEmotions taxonomy.
  - Transfer Learning: F1 score on data transferred between domain X and GoEmotions.
- Decision thresholds: No thresholds are used. The data is presented in full granularity.
- Uncertainty and variability: Repeated experiments have yielded results with similar taxonomical rankings.

The model has been evaluated on 10 publicly available datasets including 9 benchmark datasets presented in compilation by Bostan and Klinger and the GoEmotions eval set. Full details about the evaluation results can be found in the paper.

### 1.4 Data Splits

The dataset is divided into training, validating and test splits as follows:

Data Split	Number of Instances
Training	43,410
Validation	5,426
Test	5,427

Table 2: Data Split Overview

Class distributions for the training and test splits are shown in the figure 1.4.

	Emotion	Train Count	Test Count
0	admiration	4130	504
1	amusement	2328	264
2	anger	1567	198
3	annoyance	2470	320
4	approval	2939	351
5	caring	1087	135
6	confusion	1368	153
7	curiosity	2191	284
8	desire	641	83
9	disappointment	1269	151
10	disapproval	2022	267
11	disgust	793	123
12	embarrassment	303	37
13	excitement	853	103
14	fear	596	78
15	gratitude	2662	352
16	grief	77	6
17	joy	1452	161
18	love	2086	238
19	nervousness	164	23
20	optimism	1581	186
21	pride	111	16
22	realization	1110	145
23	relief	153	11
24	remorse	545	56
25	sadness	1326	156
26	surprise	1060	141
27	neutral	14219	1787

Figure 1: Class Distribution in Training and Test Splits

## 1.5 Additional Features

Although the comments vary in length, the maximum sequence length has been capped at 30 tokens in the training and evaluation datasets to ensure concise and focused emotional expressions. Furthermore, linguistic context consistency has been ensured by only choosing comments that are in English.

## 2 Fine-tuned Models

For this task, the RoBERTa-base model and the DistilBERT-base-uncased model have been used.

## 2.1 Model 1: RoBERTa-base

RoBERTa or "A Robustly Optimized BERT Pretraining Approach" is a transformer-based model that builds on the BERT architecture. It improves upon BERT's pretraining methodology by training longer with larger batches over more data, removing the next sentence prediction objective, and dynamically changing the masking pattern applied to the training data. The model has approximately 125 million total parameters.

### 2.1.1 Pretraining Dataset

RoBERTa was pretrained on a diverse and extensive corpus totalling around 160GB of text, including:

- BookCorpus which is a dataset of 11,038 unpublished books
- English Wikipedia excluding lists, tables and headers
- CommonCrawl News with over 63 million English news articles from 2016-2019
- OpenWebText which is an open source recreation the WebText dataset used to train GPT-2
- Stories which is a dataset containing a subset of CommonCrawl data filtered to match story-like style of Winograd schemas

### 2.1.2 Compute Requirements for Pretraining

The pretraining of RoBERTa-base involved substantial computational resources, including:

- **Hardware:** 1,024 V100 GPUs
- **Training Duration:** 500,000 steps
- **Batch Size:** 8,000
- **Sequence Length:** 512 tokens
- **Optimizer:** Adam with a learning rate of  $6e-4$

## 2.2 Model 2: DistilBERT-base-uncased

DistilBERT is a distilled version of BERT, designed to be smaller, faster, and more efficient while retaining 97% of BERT's language understanding capabilities. It achieves this through knowledge distillation during the pretraining phase, effectively reducing the model size by 40%. The model has approximately 66 million total parameters.

### 2.2.1 Pretraining Dataset

DistilBERT was pretrained on the same corpus as BERT, which includes:

- BookCorpus which is a dataset of 11,038 unpublished books
- English Wikipedia excluding lists, tables and headers

### 2.2.2 Compute Requirements for Pretraining

The pretraining of DistilBERT-base-uncased utilized:

- **Hardware:** 8 16GB V100 GPUs
- **Training Duration:** 90 hours

## 2.3 Fine-Tuning Steps

The RoBERT-a and DistilBERT-base-uncased models were fine-tuned on the GoEmotions dataset using the following process:

1. Dataset Preparation
  - Loaded the GoEmotions dataset with simplified labels (27 categories in total).
  - Verified the dataset structure and label distribution.
  - Used the predefined train, validation and test splits.
2. Preprocessing
  - Tokenize the data using RoBERTa’s byte-level Byte-Pair Encoding (BPE) tokenizer and DistilBERT’s WordPiece tokenizer.
  - Applied truncation and padding to a maximum sequence length of 128 tokens.
  - Converted datasets to PyTorch format.
3. Hyperparameter Optimization
  - Optimized for weighted F1 score on the validation set.
  - Tuned parameters such as learning rate, weight decay and batch size.
  - Performed 5 trials per model to balance optimization quality with computational time.
4. Model Configuration
  - Configured both models with a classification head for 27 labels.
  - Implemented early stopping with a patience of 3 epochs.
  - Used mixed precision training for memory efficiency.
  - Applied gradient accumulation (steps=2) to stimulate larger batch sizes.
5. Fine-tuning Process
  - Trained both models using the Trainer API from HuggingFace Transformers.
  - Applied the best hyperparameters found during optimization.
  - Trained for up to 5 epochs with early stopping.
  - Evaluated performance after each epoch on the validation set.
  - Saved the best checkpoint based on validation performance.